



NN Interpretability

Aleksandar Aleksandrov
Hans Hao-Hsun Hsu



Outline

- Introduction
- Model-based Methods
 - Activation Maximization
- Decision-based Methods
 - Deconvolutional Network
 - Occlusion Sensitivity
 - Saliency Maps
 - Layer-wise Relevance Propagation


Introduction

- What is NN interpretability?
 - NN interpretability refers to the process of **mapping of abstract concepts in a human-understandable domain**. A collection of features in the human-interpretable domain allows us to **provide possible explanations for the decisions of a model**.
- What types of NN interpretability methods are there?
 - **Model-based methods** (e.g. Activation Maximization) try to explain what does the concepts learned from a model look like. (How does a “dog” typically look like?)
 - **Decision-based methods** (e.g. Layerwise Relevance Propagation) try to explain why did the model assign a certain concept to a premeditated input. (Why is this example classified as “dog”?)


Activation Maximization (AM)

- AM is a **model-based approach** that searches for an **input pattern** which **elicits a maximum model response** for a class of interest.
- Variations:
 - General AM $\max_{\mathbf{x}} \log p(\omega_c | \mathbf{x}) - \lambda \|\mathbf{x}\|^2.$
 - AM with an Expert $\max_{\mathbf{x}} \log p(\omega_c | \mathbf{x}) + \log p(\mathbf{x}).$
 - AM in Code Space $\max_{\mathbf{z} \in \mathcal{Z}} \log p(\omega_c | g(\mathbf{z})) - \lambda \|\mathbf{z}\|^2,$


- Random noise (Class 6)



- Random image (Class 5)




- Mean image for class (Class 9)



Deconvolutional Network (DeConvNet)


- DeConvNet is a **decision-based approach** for mapping feature activities back to the input pixel space.
- DeConvNet has the **reversed structure** of a concrete CNN model and **reuses the initially learned weights**.
- Variations
 - DeConvNet with all filters
 - DeConvNet with a single filter in a given layer

- DeConvNet with all filters




- DeConvNet with a single filter in a given layer


Input Image:



Results for the first CONV layer

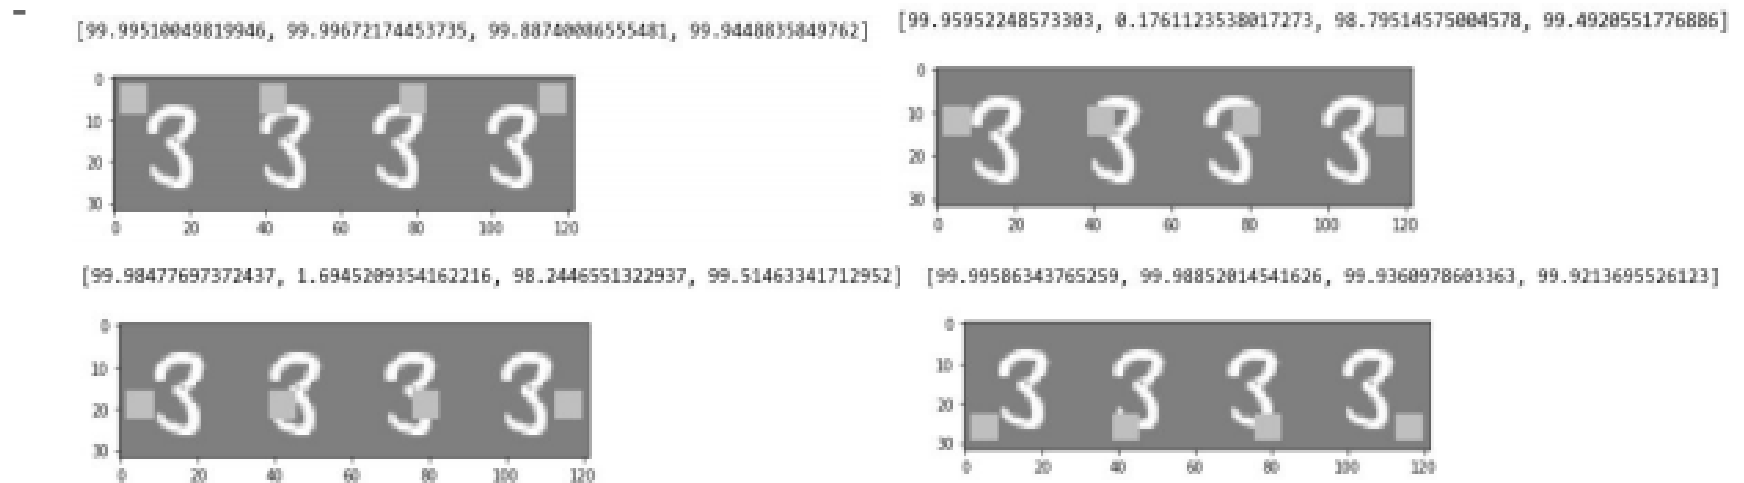


Results for the second CONV layer




Occlusion Sensitivity

- Occlusion Sensitivity is a **decision-based approach** in which parts of the input are **deliberately obstructed** to mislead the decision of the model



Saliency Maps

- Saliency Map is a **decision-based approach** that indicates which pixels need to be changed the least to affect the class score the most
- **Problem:** Doesn't highlight which pixel causes the prediction of "4"
Not conservative $f(x) = \sum_{i=1}^V R(x_i)$



Layer-wise Relevance Propagation

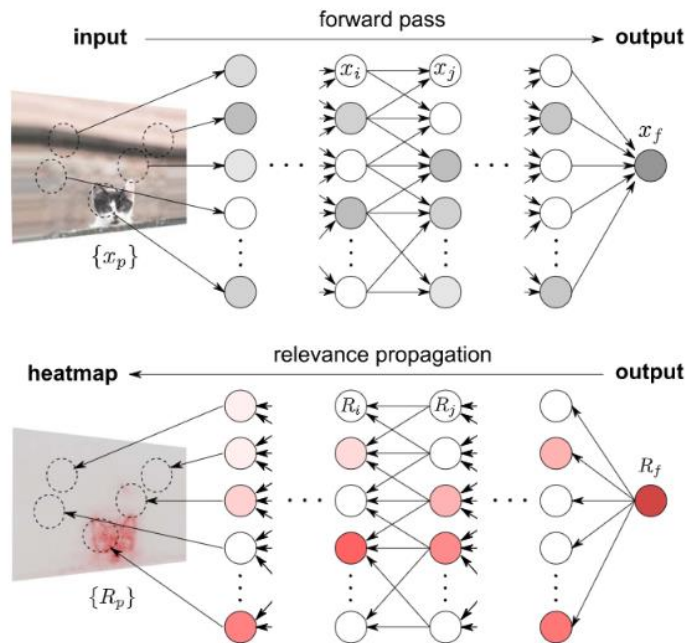
- Layer-wise Relevance Propagation(LRP) is a **decision-based approach** by propagating **relevance scores backward** using a set of purposely designed rules

- **Variations:**

- Naive LRP Rule
$$R_j = \sum_k \frac{a_j w_{jk}}{\sum_{0,j} a_j w_{jk}} R_k$$


- LRP- ϵ Rule
$$R_j = \sum_k \frac{a_j w_{jk}}{\epsilon + \sum_{0,j} a_j w_{jk}} R_k$$

- LRP- γ Rule
$$R_j = \sum_k \frac{a_j \cdot (w_{jk} + \gamma w_{jk}^+)}{\sum_{0,j} a_j \cdot (w_{jk} + \gamma w_{jk}^+)} R_k$$





Layer-wise Relevance Propagation

- LRP- ϵ Rule



- Red pixels: Raise the probability for the class "4"
- Blue pixels: Lower the probability for the class "4"
- Combine with occlusion:



Label 1; Predicted 6

Next Step

- Implementation of further methods
 - Guided Backpropagation
 - Class Activation Maps
 - AM in CodeSpace (e.g. with GAN)
 - DeepDream
- Introduction of uncertainty

Sources

- 1) Montavon, Grégoire, Wojciech Samek, and Klaus-Robert Müller. "Methods for Interpreting and Understanding Deep Neural Networks." Digital Signal Processing 73 (2018): 1–15. Crossref. Web.
- 2) Matthew D. Zeiler and Rob Fergus (2013). Visualizing and Understanding Convolutional NetworksCoRR, abs/1311.2901.
- 3) Layer-Wise Relevance Propagation: An Overview
- 4) Olah, et al., "Feature Visualization", Distill, 2017.
- 5) Olah, et al., "The Building Blocks of Interpretability", Distill, 2018.
- 6) Karen Simonyan, Andrea Vedaldi, & Andrew Zisserman. (2013). Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps.



Backup Slides




Saliency Map

$$S_c(I) \approx w^T I + b,$$

$$w = \left. \frac{\partial S_c}{\partial I} \right|_{I_0}$$

Saliency Map



$$f(x, y) = \begin{cases} x^2 + y^2 - 9 < 0 & \text{for class N(egative)} \\ x^2 + y^2 - 9 \geq 0 & \text{for class P(ositive)} \end{cases}$$

$$\nabla f(x, y) = [\partial f / \partial x \quad \partial f / \partial y] = [2x \quad 2y]$$

LRP Calculation

element-wise	vector form
$z_k \leftarrow \sum_j a_j w_{jk}^+$	$z \leftarrow W_+^\top \cdot a$
$s_k \leftarrow R_k / z_k$	$s \leftarrow R \oslash z$
$c_j \leftarrow \sum_k w_{jk}^+ s_k$	$c \leftarrow W_+ \cdot s$
$R_j \leftarrow a_j c_j$	$R \leftarrow a \odot c$

```
def lrp(layer,a,R):  
    clone = layer.clone()  
    clone.W = maximum(0,layer.W)  
    clone.B = 0  
  
    z = clone.forward(a)  
    s = R / z  
    c = clone.backward(s)  
  
    return a * c
```

$$f(x) = \dots = \sum_{d=1}^{V(l+1)} R_d^{(l+1)} = \sum_{d=1}^{V(l)} R_d^{(l)} = \dots = \sum_{i=1}^{V(1)} R_d^{(1)}$$

$$c_j = [\nabla(\sum_k z_k(\mathbf{a}) \cdot s_k)]_j$$

Deconvolutional Network

